

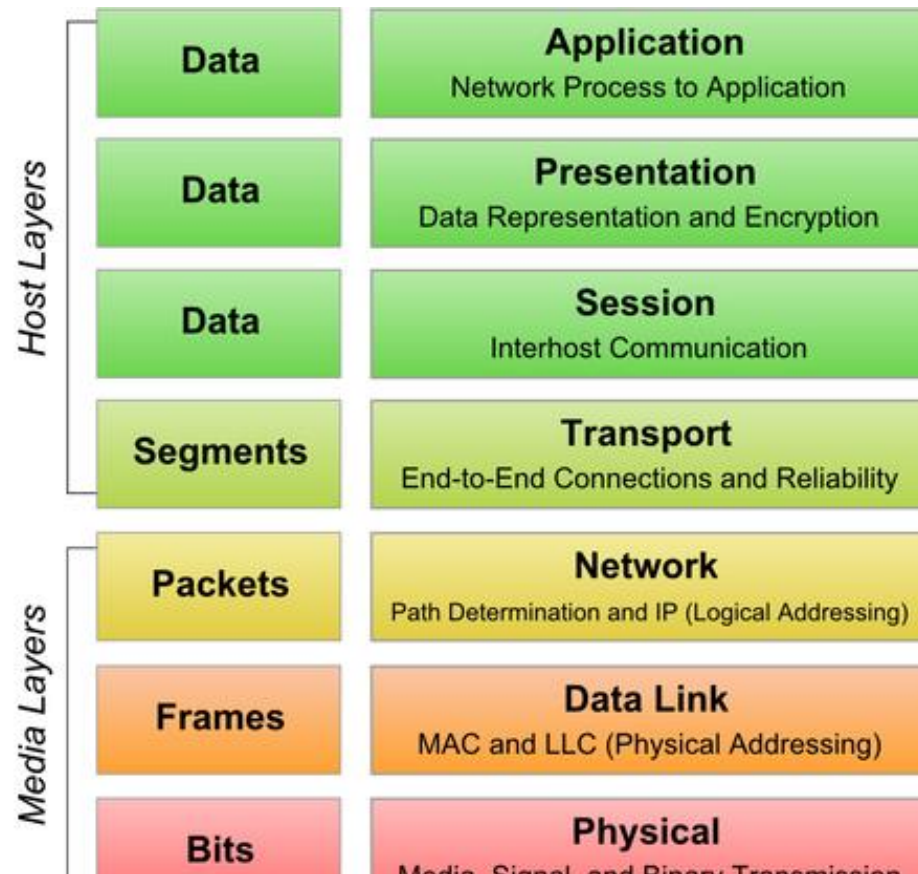
RAWDATA SECTION 2

Troels Andreassen & Henrik Bulskov




NETWORKS

- Overview
- IP
- TCP
- HTTP
- C# Networking (and Threading briefly)



LAYERS



LAYERS

PROTOCOLS

Layer

Application

Network Process to Application

Transport

End-to-End Connections and Reliability

Network

Path Determination and IP (Logical Addressing)

Response Headers

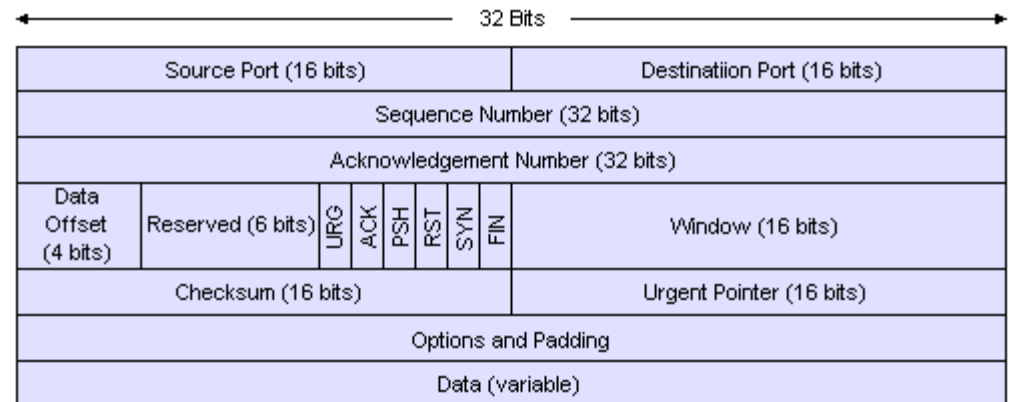
view parsed

HTTP/1.1 200 OK

← response starting line

response headers

Content-Encoding: gzip
Vary: Accept-Encoding
Transfer-Encoding: chunked
Date: Wed, 06 Mar 2013 13:25:52 GMT
Server: LiteSpeed
Connection: close
X-Powered-By: PHP/5.3.17
Expires: Wed, 11 Jan 1984 05:00:00 GMT
Last-Modified: Wed, 06 Mar 2013 13:25:52 GMT
Cache-Control: no-cache, must-revalidate, max-age=0
Pragma: no-cache
X-Frame-Options: SAMEORIGIN
Content-Type: text/html; charset=UTF-8





0	4	8	16	19	31
Version	IHL	Type of Service	Total Length		
Identification			Flags	Fragment Offset	
Time To Live		Protocol	Header Checksum		
Source IP Address					
Destination IP Address					
Options				Padding	

CORE CONCEPTS

- Bandwidth and Throughput
 - Data rate (the rate at which bits are transmitted)
 - Throughput (overall effective transmission rate)
 - Bandwidth ?
 - Goodput (the amount of usable data delivered to the receiving application)

CORE CONCEPTS

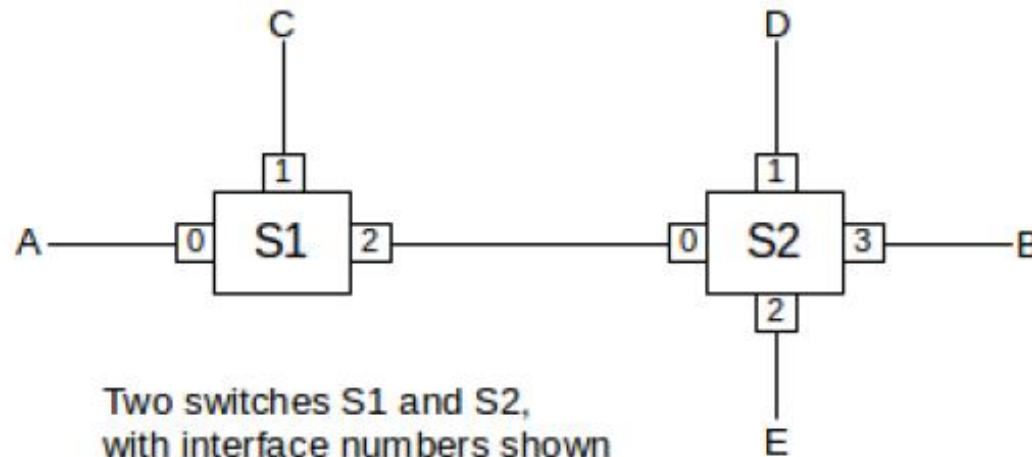
- Packets
 - Packets are modest-sized buffers of data, transmitted as a unit through some shared set of links



Single and multiple headers

CORE CONCEPTS

- Routing and Switching
 - In the datagram-forwarding model of packet delivery, packet headers contain a destination address. It is up to the intervening switches or routers to look at this address and get the packet to the correct destination




CORE CONCEPTS

- Congestion
 - Problem
 - packets arriving faster than they can be sent out
 - multiple inputs and all destined for the same output
 - Solution
 - queue incoming packets
 - drop packets

CORE CONCEPTS

- LAN
 - physical links that are, ultimately, serial lines
 - common interfacing hardware connecting the hosts to the links
 - protocols to make everything work together
- Ethernet
 - hardware address or MAC (Media Access Control) address
 - broadcast
 - unicast
 - switched ethernet

THE INTERNET



WHAT IS A PROTOCOL?

- In diplomatic circles, a protocol is the set of rules governing a conversation between people
- The client and server carry on a machine-to-machine conversation
- A network protocol is the set of rules governing a conversation between a client and a server

NETWORK PROTOCOLS

- The details are only important to developers.
- The rules are defined by the inventor of the protocol – may be a group or a single person.
- The rules must be precise and complete so programmers can write programs that work with other programs.
- The rules are often published as an RFC along with running client and server programs.

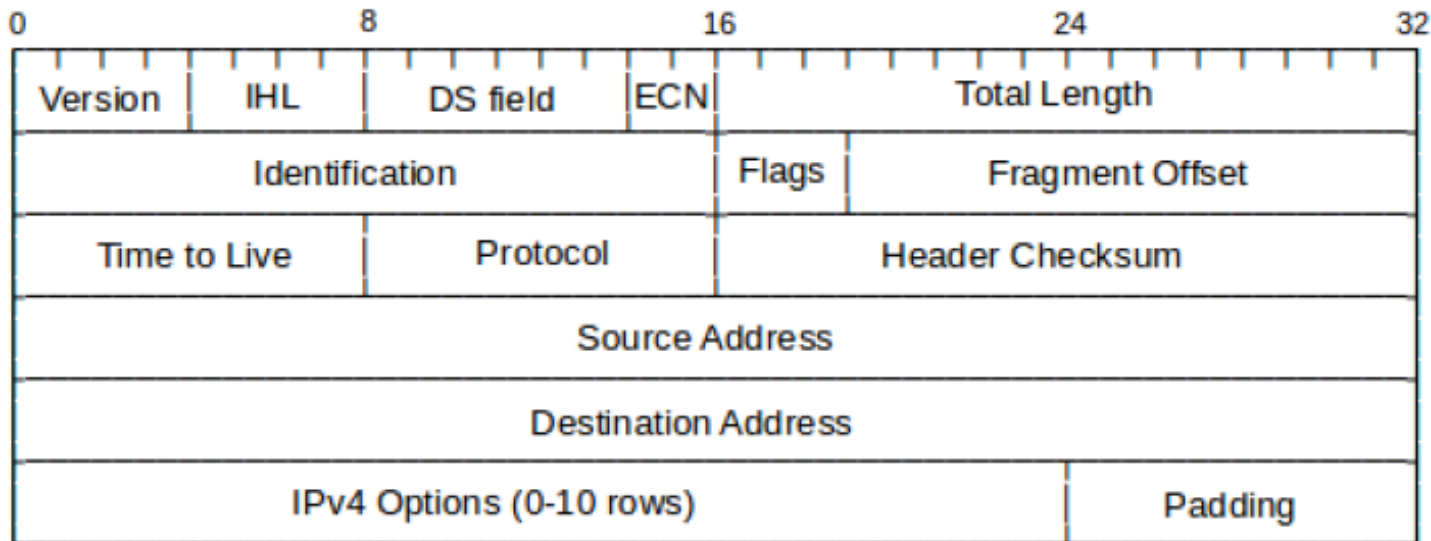
RFC = request for comments

IP PROTOCOL

- point-to-point links
- support universal connectivity
(everyone can connect to everyone else)
- IPv4 (4 bytes)
- IPv6 (16 bytes)

IP PROTOCOL

- destination and source addresses
- indication of ipv4 versus ipv6
- a Time To Live (TTL) value, to prevent infinite routing loops
- a field indicating what comes next in the packet (e.g. TCP v UDP)



IP ADDRESS

- IP provides a global mechanism for addressing and routing
- An essential feature of IPv4 (and IPv6) addresses is that they can be divided into a “network” part (a prefix) and a “host” part (the remainder).

first few bits	first byte	network bits	host bits	name	application
0	0-127	8	24	class A	a few very large networks
10	128-191	16	16	class B	institution-sized networks
110	192-223	24	8	class C	sized for smaller entities

- IP addresses, unlike Ethernet addresses, are administratively assigned

IP ADDRESS (IPV4)

- A 32 bit address that is used to uniquely identify a computer on a network
- The Network ID portion of the IP Address identifies the network where the computer sits
- The Host ID portion of the IP Address uniquely identifies the computer on its network

IP Address:	192.168.10.1	192.168.10.1	192.168.10.1
Subnet Mask:	255.255.255.0	255.255.0.0	255.0.0.0
Addresses:	254	65,534	16,777,214
Class:	C	B	A
CIDR:	192.168.10.1/24	192.168.10.1/16	192.168.10.1/8

SUBNETS



11000000.10101000.00001010.00000000	192.168.10.0
11000000.10101000.00001010.00000001	192.168.10.1
11000000.10101000.00001010.00000010	192.168.10.2
11000000.10101000.00001010.00000011	192.168.10.3
11000000.10101000.00001010.00000100	192.168.10.4
11000000.10101000.00001010.00000101	192.168.10.5
11000000.10101000.00001010.00000110	192.168.10.6
11000000.10101000.00001010.00000111	192.168.10.7

Mask: 255.255.255.248

SUBNETS

11000000.10101000.00001010.00000000	192.168.10.0
11000000.10101000.00001010.00000001	192.168.10.1
11000000.10101000.00001010.00000010	192.168.10.2
11000000.10101000.00001010.00000011	192.168.10.3
11000000.10101000.00001010.00000100	192.168.10.4
11000000.10101000.00001010.00000101	192.168.10.5
11000000.10101000.00001010.00000110	192.168.10.6
11000000.10101000.00001010.00000111	192.168.10.7

Network ID: 192.168.10.0

Host ID's: 192.168.10.1 – 192.168.10.6

Broadcast ID: 192.168.10.7

SUBNETS


11000000.10101000.00001010.00001000	192.168.10.8
11000000.10101000.00001010.00001001	192.168.10.9
11000000.10101000.00001010.00001010	192.168.10.10
11000000.10101000.00001010.00001011	192.168.10.11
11000000.10101000.00001010.00001100	192.168.10.12
11000000.10101000.00001010.00001101	192.168.10.13
11000000.10101000.00001010.00001110	192.168.10.14
11000000.10101000.00001010.00001111	192.168.10.15

Network ID: 192.168.10.8

Host ID's: 192.168.10.9 – 192.168.10.14

Broadcast ID: 192.168.10.15

- almost a null protocol



- unreliable
- common to use UDP as basis for a Remote Procedure Call
- well-suited for “request-reply” semantics
- popular for real-time transport

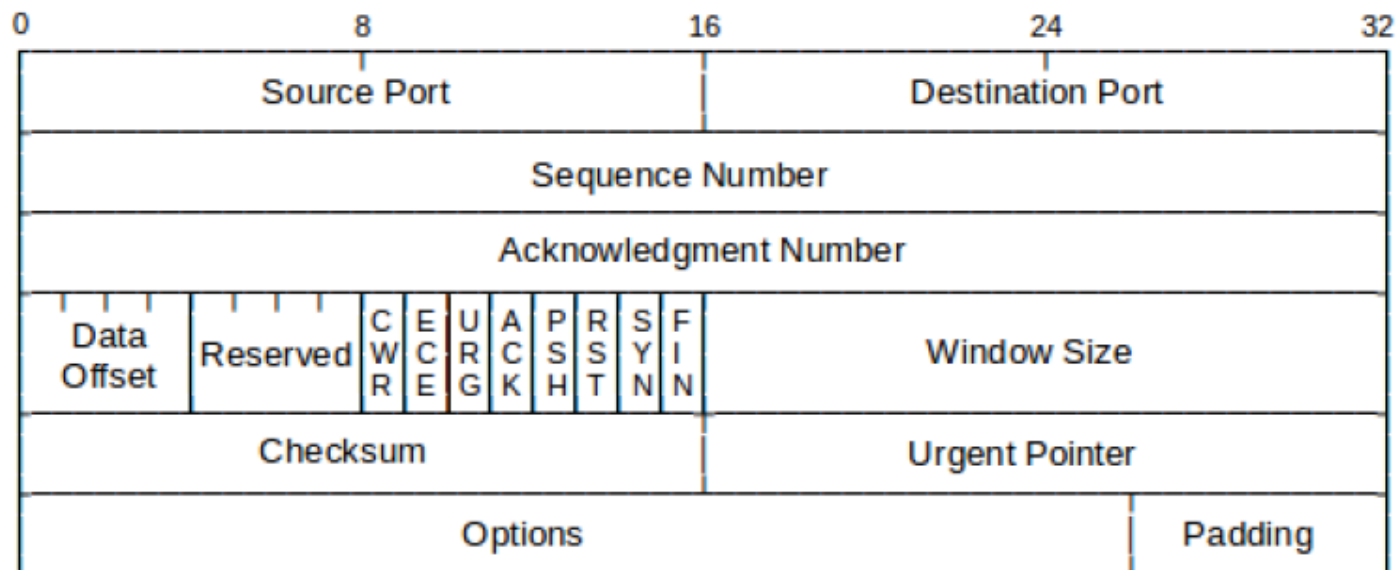
UDP PROTOCOL

TCP PROTOCOL

- **reliability:** TCP numbers each packet, and keeps track of which are lost and retransmits them after a timeout, and holds early-arriving out-of-order packets for delivery at the correct time. Every arriving data packet is acknowledged by the receiver; timeout and retransmission occurs when an acknowledgment isn't received by the sender within a given time.
- **connection-orientation:** Once a TCP connection is made, an application sends data simply by writing to that connection. No further application-level addressing is needed.
- **stream-orientation:** The application can write 1 byte at a time, or 100KB at a time; TCP will buffer and/or divide up the data into appropriate sized packets.
- **port numbers:** these provide a way to specify the receiving application for the data, and also to identify the sending application.
- **throughput management:** TCP attempts to maximize throughput, while at the same time not contributing unnecessarily to network congestion.

TCP HEADER

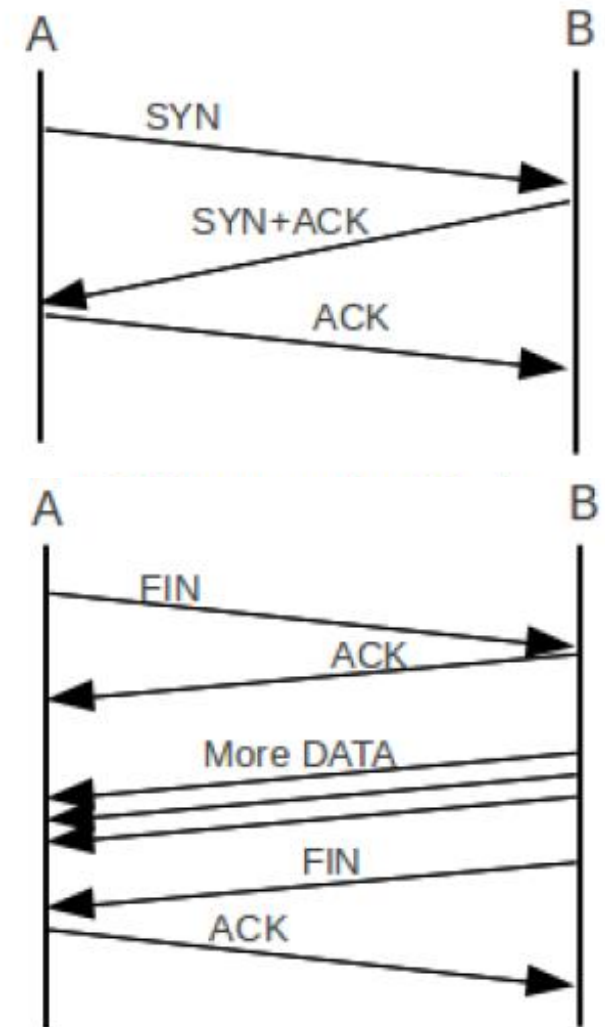
- sequence and acknowledgment numbers are for numbering the data



- flags:
 - SYN: for SYNchronize; marks packets that are part of the new-connection handshake
 - ACK: indicates that the header Acknowledgment field is valid; that is, all but the first packet

TCP CONNECTION ESTABLISHMENT

- The handshake proceeds as follows
 - A sends B a packet with the SYN bit set (a SYN packet)
 - B responds with a SYN packet of its own; the ACK bit is now also set
 - A responds to B's SYN with its own ACK
- Close the connection
 - A sends B a packet with the FIN bit set (a FIN packet), announcing that it has finished sending
 - B sends A an ACK of the FIN
 - When B is also ready to cease sending, it sends its own FIN to A
 - A sends B an ACK of the FIN; this is the final packet in the exchange



HYPertext TRANSfer PROTOCOL

- A high level protocol built on top of a TCP connection for exchanging messages (with *arbitrary content*)
 - Each (request) message from client to server is followed by a (response) message from server to client.
 - Facilitates the remote invocation of methods on the server.
- Web: A set of client and server processes on the Internet that communicate via HTTP.

Application
TCP
IP
Link Level Protocol

http

←----- Added features to support client interactions (reliability flow control, ..)

←----- End-to-end protocol

←----- Protocol for transmitting packets between neighboring nodes

PROTOCOL STACK

CLIENTS AND SERVERS

- Client: browser capable of displaying HTML pages.
- Web Server: stores pages for distribution to clients.
- Pages identified by Uniform Resource Locator (URL).
 - *<protocol>: protocol to be used to communicate with host.*
- Example - http, ftp

Start line: *<method> <URL>*
 <protocol_version> CrLf
Followed by: *<header>**
Followed by: CrLf
Followed by: *<data>*

there can
be several
header lines

<method> = GET | HEAD | POST | PUT |
<protocol_version> = HTTP/1.1 |

HTTP REQUEST FORMAT

REQUEST METHODS

- **GET** – response body contains data identified by argument URL
- **HEAD** – response header describes data identified by argument URL (no response body)
 - Use: has page changed since last fetched?
- **PUT** – request body contains page to be stored at argument URL
- **DELETE** – delete data at argument URL
- **POST** – request body contains a new object to be placed subordinate to object at argument URL
 - Use: adding file to directory named by URL
 - Use: information entered by user on displayed forms
- Others

CLIENT/SERVER INTERACTION

1. User supplies URL (clicks on link)
`http://yourbusiness.com/~items/printers.html`
2. Browser translates *<host_name>* (*yourbusiness.com*) to *host_ip_address*
3. Browser assumes a port number of 80 for http (if no port is explicitly provided as part of *<host_name>*)
 - Program at port 80 interprets http headers
4. Browser sets up TCP connection to yourbusiness.com at (*host internet address, 80*)
5. Browser sends http message
`GET ~items/printers.html HTTP/1.0`

HTTP RESPONSE

Status line: *<HTTP_version>*
<status_code>
<reason_line> CrLf

Followed by: *< header >**

Followed by: *<data>*

HTTP RESPONSE

<status_code> = 3 digits

Ex: 2xx -- success

4xx -- bad request from client

5xx -- server failed to fulfill valid request

<reason_line> = explanation for human reader

<header> = <field_name> : <value> CrLf

<field_name> =

Allowed | -- methods supported by URL

Date | -- creation date for response

Expires | -- expiration date for data

Last-Modified | -- creation date for object

Content-Length | Content-Type |

CLIENT/SERVER INTERACTION

6. Server sends response message with requested html page to browser

HTTP/1.0 200 Document follows

Date: *<date>*

Content-Type: text/html

Content-Length: *integer*

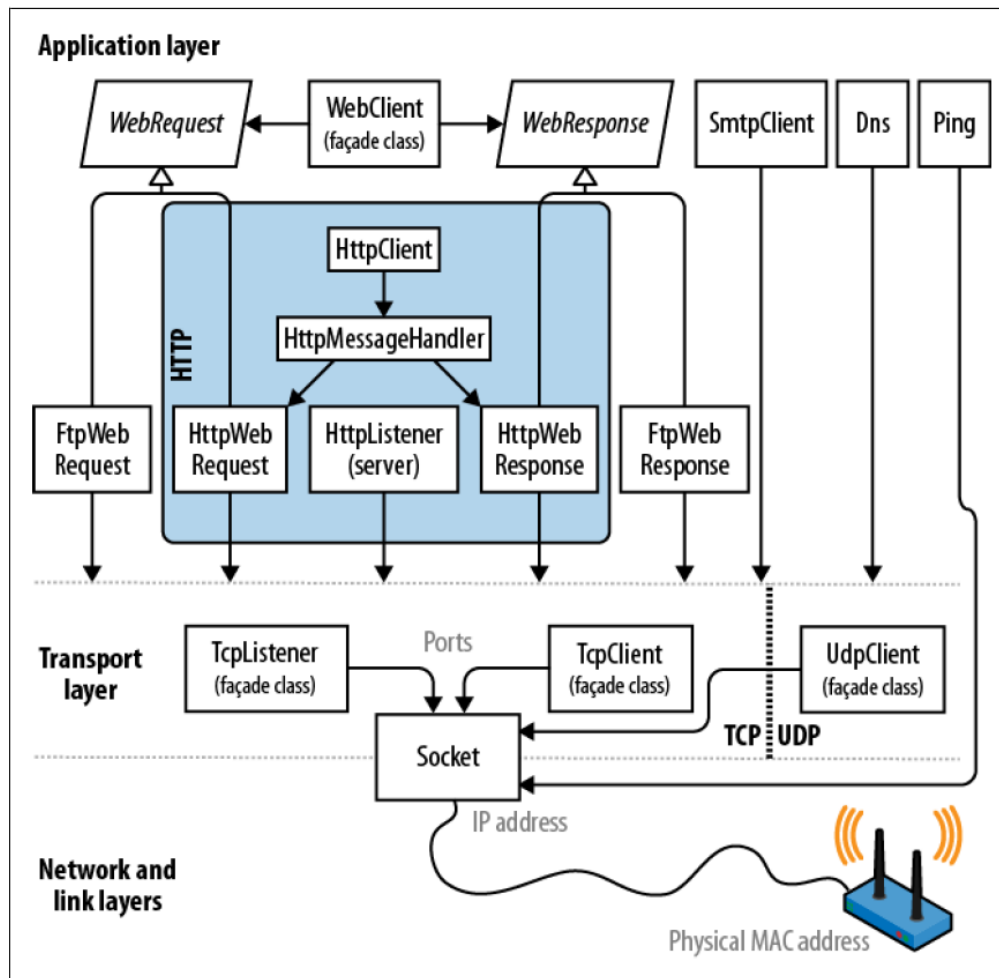
Expires: *date*

html document ~items/printers.html goes here

7. Server releases TCP connection (stateless)
8. Browser receives page and displays it

HTTP PERFORMANCE

- HTTP/1.0 allowed one transaction per connection
 - TCP connection setup and teardown are expensive
 - TCP's *slow start* slows down the initial phase of data transfer
 - typical Web pages use between 10-20 resources (HTML + images)
 - typically, these resources are stored on the same server
- HTTP/1.1 introduces *persistent connections*
 - the TCP connection stays open for some time (10 sec is a popular choice)
 - additional requests to the same server use the same TCP connection
- HTTP/1.1 introduces *pipelined connections*
 - instead of waiting for a response, requests can be queued
 - the server responds as fast as possible
 - the order may not be changed (there is no sequence number)



C# NETWORKING

Sockets

TcpClient/TcpListener

HttpClient

C# NETWORKING

SOCKETS

```
var host = "roskilde.dk";  
var ipHostEntry = Dns.GetHostEntry(host);  
var ipAddress = ipHostEntry.AddressList[0];  
var ipEndPoint = new IPEndPoint(ipAddress, 80);  
var socket = new Socket(  
    AddressFamily.InterNetwork, SocketType.Stream, ProtocolType.Tcp);
```

TCPCLIENT/TCPLISTENER

```
using (TcpClient client = new TcpClient())
{
    client.Connect ("address", port);
    using (NetworkStream n = client.GetStream())
    {
        // Read and write to the network stream...
    }
}

TcpListener listener = new TcpListener (<ip address>, port);
listener.Start();

while (keepProcessingRequests)
    using (TcpClient c = listener.AcceptTcpClient())
        using (NetworkStream n = c.GetStream())
        {
            // Read and write to the network stream...
        }

listener.Stop();
```

CONCURRENCY AND ASYNCHRONY

- Threads

- A *thread* is an execution path that can proceed independently of others

```
Thread t = new Thread (WriteY);  
t.Start();  
  
static void WriteY()  
{  
    for (int i = 0; i < 1000; i++) Console.Write ("y");  
}
```

- Join and Sleep

- You can wait for another thread to end by calling its Join method
 - Thread.Sleep pauses the current thread for a specified period

- Tasks

```
Task.Run (() => Console.WriteLine ("Foo"));  
  
new Thread (() => Console.WriteLine ("Foo")).Start();
```

- Task.Delay

- Task.Delay is the *asynchronous* equivalent of Thread.Sleep

PRINCIPLES OF ASYNCHRONY

- Synchronous Versus Asynchronous Operations
 - A *synchronous operation* does its work *before* returning to the caller
 - An *asynchronous operation* does (most or all of) its work *after* returning to the caller

- Awaiting

```
var result = await expression;  
statement(s);
```

```
var awaiter = expression.GetAwaiter();  
awaiter.OnCompleted (() =>  
{  
    var result = awaiter.GetResult();  
    statement(s);  
});
```

- Async